

Polite Bursting: Feedback Admission Control under Age-Limited Sensing on Shared Clusters

Anonymous Author(s)

Paper submitted for double-blind review

Abstract—Research teams run batch GPU work on shared clusters they do not own. Fair-use policy caps how many pods a guest may hold and requires each pod to stay busy, so a guest has to restrain itself. How much capacity a guest can actually use is set by the other tenants. It is not announced, and a guest learns it only by probing. Every probe is already out of date by the time a controller acts on it. We show that acting on stale capacity readings has an exact ceiling. For a contended resource that is either free or taken, the extra goodput any feedback controller can win over the best fixed budget equals half the correlation between capacity now and capacity as last seen. A rule that admits work whenever the last reading showed the resource free attains that ceiling. The ceiling falls to zero once the sensing delay grows past the time over which neighbour load stays correlated, so past that point no feedback controller beats a fixed budget. We turn the result into two algorithms. The first estimates the correlation from its own probe stream and switches between feedback and a fixed budget, which holds it at or above the better of the two everywhere. The second admits work across several contended units by treating each unit as a separate binary decision, which extends the result beyond a single unit. Short controlled experiments support both. On a two-GPU node driven by a real competing tenant, feedback admission under structured neighbour load delivers 1.41 times the goodput of a fixed budget while over-accepting a quarter as often as a greedy policy, and it leaves the competing tenant more of its own throughput. Under memoryless load feedback only breaks even, which is what the ceiling predicts, and the switching algorithm detects that case and gives up the feedback loop rather than paying for it.

I. INTRODUCTION

Much machine-learning work is bursty. Long stretches of local development are punctuated by short runs of batch jobs, sweeps, fine-tunes and evaluations, that briefly need far more GPUs than a team owns. Shared clusters such as the National Research Platform [1] make that capacity available on guest terms. Fair-use policy caps concurrent pods per user, requires sustained utilization of what is requested, forbids idle or sleep-only work, and reclaims resources on completion. The guest also sits outside the cluster network, often reachable only through one forwarded port or a mesh VPN, with no inbound access.

Two obvious strategies both fail here. Submitting everything at once takes more than a fair share, leaves pods pending or idle, and trips the thresholds that get a namespace banned. A fixed submission rate is safe but wastes the guest's own throughput whenever capacity is free. Neither one reacts to what the cluster is doing.

Being a good guest is a control problem, and how well it can be solved depends on how stale the guest's view of the

cluster is. This paper makes that dependence exact and then builds admission algorithms on it.

Our first result is a ceiling on the value of feedback. When a controller acts on a capacity reading of a given age, the extra goodput it can win over the best fixed budget equals half the correlation between capacity now and capacity as last seen. The ceiling is tight, since a rule that admits work whenever the last reading showed capacity free reaches it. It also implies an impossibility. Once the sensing delay exceeds the time over which neighbour load stays correlated, the correlation is near zero and no feedback controller can beat a fixed budget.

Our second result turns the ceiling into algorithms. The correlation a controller needs is measurable from the probe stream the controller already collects, so a guest can estimate it online and act on the estimate. Algorithm 1 does this with one scalar of state, switching between a feedback loop and a fixed budget, and it never falls below the better of the two by more than a bounded amount. Algorithm 2 admits work across several contended units by running one binary decision per unit, which lifts the single-unit ceiling to any integer capacity.

Our third contribution is measurement rather than proof. The sensing delay is a physical quantity and we measure it at roughly 2.5 seconds on real hardware, dominated by CUDA start-up rather than by polling. We then drive a two-GPU node with a real competing tenant and show the predicted behaviour on both sides of the crossover, including a case where feedback stops paying for itself and the switching algorithm gives it up.

The control plane runs inside the cluster and reaches the home node over a single outbound port. The fabric that carries it is characterized in a companion paper [2] and is treated here as given, since what matters for the result is only that the guest's view of capacity arrives late.

II. BACKGROUND AND MOTIVATION

Host policy. The reference host [1] publishes a fair-use policy that our design treats as ground truth. A user may hold at most $K=4$ concurrent pods. Per running pod, utilization relative to request must stay in published bands, at least 40% for GPU and 20 to 200% for CPU. Jobs must do real work, since sleep-only jobs are banned, and must release resources on completion. Idle storage volumes are reclaimed. Administrators monitor utilization and ban namespaces that hold under-utilized requests, so politeness is not optional.

Network reality. The guest reaches the cluster over one forwarded port or a mesh VPN, with no inbound access. A control plane must therefore live behind that boundary and




Fig. 1. PoliteBurst architecture. The home client submits over a federated NATS leaf link. The in-cluster controller admits jobs, pacing concurrency to the politeness budget, and launches GPU worker pods whose results route back through the leaf. The dashed arc is the stale feedback the controller acts on, and its delay bounds the achievable advantage by Thm. 1.

communicate over a fabric that traverses it, which motivates the leaf-node design of §VI.

Why existing tools do not settle it. Workflow and function-execution frameworks such as HTCondor [3], funcX and Globus Compute [4], Parsl [5], Ray [6] and Dask [7] scale workers to demand and carry no notion of a host politeness budget. Left to default scaling they over-provision under contention. We make self-restraint a controlled objective instead.

III. SYSTEM MODEL

A guest bursts batch GPU pods onto a shared cluster under a fair-use policy that allows at most K concurrent pods and requires each admitted pod to sustain a utilization floor. Time is slotted into control epochs $t \in \mathbb{N}$. Let $a_t \in \{0, \dots, C\}$ be the capacity the cluster can productively give the guest at epoch t . It is set by the other tenants, so it is exogenous, stochastic and not directly observable. The guest chooses a window $w_t \in \{0, \dots, K\}$. Useful goodput is $G_t = \min(w_t, a_t)$, and the surplus $(w_t - a_t)^+$ is over-admitted, since those pods sit below the utilization floor and mark the guest impolite. Writing $\rho \triangleq \limsup_T \frac{1}{T} \sum_{t < T} \mathbf{1}[w_t > a_t]$ for the long-run over-admission fraction, the guest solves

$$\max_{\{w_t\}} \mathbb{E}[\min(w_t, a_t)] \quad \text{s.t.} \quad \underbrace{w_t \leq K \forall t}_{\text{hard cap}}, \quad \underbrace{\rho \leq \epsilon}_{\text{policy}}. \quad (1)$$

The constraint is two-sided, since the guest must stay below the cap and above the floor, which already separates the problem from loss-only congestion control.

Age-limited sensing. The guest does not observe a_t . It probes cluster state and the probe is stale. We model this as an age of information D , meaning that at epoch t the controller's freshest reading of capacity has age D , so admissible policies are causal functions of $\{a_{s-D}\}_{s \leq t}$. Here D is a physical and measured quantity. Sec. VIII puts it near 2.5 s, dominated by CUDA start-up rather than by the 27 ms poll. A controller can only act on a D -old picture of a cluster the neighbours are still changing.

The contended unit. Politeness is decided at the margin, so we fix attention on the single contended resource unit whose availability fluctuates with the neighbour. Let $x_t \in \{0, 1\}$

indicate that this unit is free ($x_t = 1$) or taken ($x_t = 0$), so admitting it when taken is the over-admission the policy punishes. We take x_t stationary with $\Pr[x_t = 1] = \frac{1}{2}$, a fully contended unit, and note that the asymmetric case only rescales the constants below. The guest's marginal decision $m_t \in \{0, 1\}$ is a possibly randomized causal function of the age- D observation $\hat{x}_t \triangleq x_{t-D}$. Marginal goodput and over-admission are

$$g \triangleq \Pr[m_t = 1, x_t = 1], \quad \rho \triangleq \Pr[m_t = 1, x_t = 0]. \quad (2)$$

Everything below concerns the achievable (ρ, g) region under age D . The stationary ρ of (2) is the ergodic limit of the epoch fraction in (1), so we reuse the symbol, and write $\bar{\rho}$ for the surplus amount in Prop. 2.

Proposition 1 (Hard-cap safety). *Any window update that clamps additive increase at K and never increases w_t on a congestion signal keeps $w_t \leq K$ for every t deterministically, so the pod cap is never violated.*

Proof. Induction on t . Increase is clamped to K and decrease cannot exceed the current value. $\square$

IV. HOW MUCH STALE FEEDBACK IS WORTH

We now characterize exactly how much closed-loop admission gains over a fixed budget when the only signal is a D -old observation. The answer is one interpretable quantity, the autocorrelation of capacity at the sensing lag.

Lemma 1 (Predictive disagreement). *Let $q(D) \triangleq \Pr[x_t \neq x_{t-D}]$ be the lag- D disagreement probability and $r(D) \triangleq \mathbb{E}[(2x_t - 1)(2x_{t-D} - 1)]$ the lag- D autocorrelation of centered availability. Then $r(D) = 1 - 2q(D) \in [-1, 1]$, and by stationarity and symmetry the joint law of $(\hat{x}_t, x_t)$ is $\Pr[\hat{x}_t = i, x_t = j] = \frac{1}{4}(1 + (2i - 1)(2j - 1)r(D))$. For a symmetric two-state Markov capacity with per-epoch flip probability $p \in (0, \frac{1}{2})$ we have $r(D) = (1 - 2p)^D = e^{-D/\tau_c}$ with coherence time $\tau_c \triangleq -1/\ln(1 - 2p) \approx 1/(2p)$.*

Proof. $r = \Pr[\text{agree}] - \Pr[\text{disagree}] = 1 - 2q$. The joint law follows from the two symmetric marginals and the single correlation. For the Markov chain, writing $\lambda_2 \triangleq 1 - 2p$ for the second eigenvalue, the D -step transition is $P^D = \frac{1}{2} \begin{pmatrix} 1 + \lambda_2^D & 1 - \lambda_2^D \\ 1 - \lambda_2^D & 1 + \lambda_2^D \end{pmatrix}$, so $\Pr[\text{disagree}] = \frac{1}{2}(1 - \lambda_2^D)$ and $r = \lambda_2^D$. The two-state chain's autocorrelation and its total-variation distance to stationarity are standard mixing facts [8]. $\square$

Theorem 1 (Achievable region and the feedback advantage). *For the symmetric two-state Markov capacity of Lemma 1, under age- D sensing the achievable (ρ, g) region of (2) is the quadrilateral with vertices $\{(0, 0), (\frac{q}{2}, \frac{1-q}{2}), (\frac{1}{2}, \frac{1}{2}), (\frac{1-q}{2}, \frac{q}{2})\}$ with $q = q(D)$, whose efficient upper boundary is the polyline $(0, 0) \rightarrow (\frac{q}{2}, \frac{1-q}{2}) \rightarrow (\frac{1}{2}, \frac{1}{2})$. The best observation-oblivious policy is confined to the diagonal $g = \rho$. Consequently the goodput gain of feedback over a fixed budget, maximized over matched over-admission levels, is*

$$\Delta(D) = \frac{1}{2} r(D) = \frac{1}{2} (1 - 2q(D)) \quad (3)$$

attained at the kink $\rho = q/2$ by the threshold policy $m_t = \hat{x}_t$, while at a lower level $\rho < q/2$ the gain is the smaller $\rho(1 - 2q)/q$. The peak value of closed-loop admission is therefore exactly half the autocorrelation of available capacity at the sensing lag (Fig. 2a). Optimality of memoryless threshold policies for symmetric two-state Markov observation is classical in opportunistic access [9], and what (3) adds is the exact value of the resulting advantage over the best fixed budget.

Proof. For the symmetric two-state Markov chain the freshest sample $\hat{x}_t = x_{t-D}$ is a sufficient statistic for x_t given the whole delayed history $\{x_{s-D}\}_{s \leq t}$, since the Markov property makes older samples conditionally uninformative. An optimal causal age- D policy therefore depends on the history only through $\hat{x}_t$, so we may restrict without loss to policies parameterized by a pair (θ_1, θ_0) with $\theta_i = \Pr[m_t=1 \mid \hat{x}_t=i]$. Using Lemma 1, $g = \frac{1}{2}(\theta_1(1-q) + \theta_0q)$ and $\rho = \frac{1}{2}(\theta_1q + \theta_0(1-q))$. This linear map sends the square $[0, 1]^2$ to the quadrilateral with the stated vertices, whose upper boundary is the stated hull. Maximizing $g - \lambda\rho$ turns on θ_1 before θ_0 because $q < \frac{1}{2}$ makes the free-observation state strictly more goodput-efficient, so the maximum-goodput point for $\rho \leq q/2$ is the kink $(\frac{q}{2}, \frac{1-q}{2})$, achieved by $m_t = \hat{x}_t$. A fixed policy ignores $\hat{x}_t$, so $\theta_1 = \theta_0 = \theta$ and $g = \rho = \theta/2$, the diagonal. The vertical distance at $\rho = q/2$ is $\frac{1-q}{2} - \frac{q}{2} = \frac{1-2q}{2} = \frac{1}{2}r(D)$. $\square$

Corollary 1 (Impossibility as sensing ages). $\Delta(D) \rightarrow 0$ as $r(D) \rightarrow 0$. For the symmetric two-state Markov model $\Delta(D) = \frac{1}{2}e^{-D/\tau_c}$, so no causal controller with sensing age D outperforms a fixed budget once $D \gg \tau_c$. Closed-loop admission helps only to the extent that the guest senses the neighbour faster than the neighbour changes. In this model low autocorrelation and unpredictability coincide, but in general they need not, since a decorrelated process with $r(D)=0$ can still be exactly predictable from its history, as a random-phase deterministic square wave is. For an arbitrary capacity process the vanishing of the advantage is therefore governed by the β -mixing coefficient rather than by $r(D)$. At $D = 0$ we have $\Delta = \frac{1}{2}$, and the advantage halves for every $\tau_c \ln 2$ epochs of added age (Fig. 2b). App. A makes the general statement precise and extends the impossibility to arbitrary non-binary and non-Markov capacity processes.

Provenance and scope. The ingredients of (3) are classical and we claim none of them. The value of information is nonnegative and convex in the prior [10], it is bounded by payoff oscillation times the total-variation movement of beliefs [11], that movement at sensing age D is the β -mixing coefficient [12], and for the symmetric two-state chain the mixing facts and threshold optimality are standard [8], [9]. What (3) contributes is their assembly under age-limited exogenous sensing, giving an exact identity with threshold achievability, a mixing converse in App. A, and confirmation on real cluster contention in §VIII. The exogenous structure matters. Unlike the driven-plant setting of networked control, the admission

action here does not affect capacity, and that is what collapses the value of feedback onto a single predictability coefficient.

Corollary 2 (The two contention profiles are two points on (3)). The contention study of §VIII instantiates this two-state $C=2$ model. Structured square load has $\tau_c \gg D$, so $r(D)$ is near 1 and feedback is near-maximally useful. Memoryless Poisson load drives $\tau_c \lesssim D$, so $r(D)$ is near 0 and feedback collapses onto the fixed budget, which is the pre-registered null. The single curve (3) predicts both cases and the crossover between them.

The binary reduction is not a special case but the atom of the general problem.

Proposition 2 (Layer-cake reduction for $C > 2$). Let $a_t \in \{0, \dots, C\}$ and score over-admission by the surplus amount $\bar{\rho} \triangleq \mathbb{E}[(w_t - a_t)^+]$. With $x_t^u \triangleq \mathbf{1}[a_t \geq u]$ and $\pi_u \triangleq \Pr[x_t^u = 1]$, both objectives decompose additively over units, since $\min(w_t, a_t) = \sum_{u \geq 1} \mathbf{1}[w_t \geq u] x_t^u$ and $(w_t - a_t)^+ = \sum_{u \geq 1} \mathbf{1}[w_t \geq u](1 - x_t^u)$. The age- D window $w_t = a_{t-D}$ maximizes the advantage layer by layer, and its total goodput advantage over the best comonotone fixed budget at matched surplus obeys

$$\Delta(D) \leq \sum_{u=1}^K \pi_u r_u(D), \quad r_u(D) = \text{corr}(x_t^u, x_{t-D}^u), \quad (4)$$

with equality if and only if $\pi_u(1 - r_u(D))$ is non-increasing in u . For a single symmetric contended unit with $\pi_u = \frac{1}{2}$ this recovers $\Delta = \frac{1}{2}r(D)$, and equality also holds whenever $r_u(D)$ does not vary with u .

Proof. The two identities are the layer-cake expansions of $\min$ and $(\cdot)^+$. Setting $w_t = a_{t-D}$ makes $\mathbf{1}[w_t \geq u] = x_{t-D}^u$, which is the Thm. 1 threshold policy applied to each layer, and the layers' nesting $x_t^1 \geq x_t^2 \geq \dots$ is inherited by the observations, so the per-unit thresholds compose into a valid scalar window. For an asymmetric layer with $\Pr[x_t^u=1] = \pi_u$, writing $P_{11}^u = \Pr[x_t^u=1, x_{t-D}^u=1]$, the same linear program as Thm. 1 gives per-unit advantage $\text{Cov}(x_t^u, x_{t-D}^u)/(1 - \pi_u) = \pi_u r_u(D)$ over the fixed policy matching that layer's surplus, and additivity of goodput and surplus sums these to the right-hand side of (4).

The bound is an inequality rather than an identity because the summed comparator need not be realizable. Matching layer u 's surplus requires the fixed policy to admit layer u with probability $\theta^u = (\pi_u - P_{11}^u)/(1 - \pi_u) = \pi_u(1 - r_u(D))$, while any scalar window satisfies $\{w_t \geq u+1\} \subseteq \{w_t \geq u\}$ and therefore forces θ^u to be non-increasing in u . Since π_u decreases in u but $1 - r_u(D)$ may increase, the per-layer matched comparator is a feasible window exactly when $\pi_u(1 - r_u(D))$ is non-increasing, which gives the stated equality condition. When it fails, no single fixed budget matches every layer's surplus at once, the true comparator is strictly better than the layerwise one, and the sum overstates the advantage. $\square$

At $C=2$ only the contended layer fluctuates, with $\pi_1=1$, $r_1=0$ and $\pi_2=\frac{1}{2}$, so (4) collapses to $\frac{1}{2}r_2(D)$. The binary law





Fig. 2. (a) Under age- D sensing the achievable (ρ, g) pairs form the shaded hull. Fixed budgets lie on the diagonal $g = \rho$, and the closed-loop optimum sits a vertical $\Delta(D) = \frac{1}{2}r(D)$ above it (Thm. 1). (b) The same value $\frac{1}{2}e^{-D/\tau_c}$ over the sensing-age and coherence-time plane with τ_c on a log scale. Points are the measured Markov contention sweep of Cor. 1, each shaded by the adaptive controller's closed-loop fraction, opening onto the fixed budget as D crosses τ_c .

is exactly the $C=2$ instance and the experiments measure it directly.

V. TWO ALGORITHMS

Theorem 1 leaves a shortfall that no fixed controller closes. Additive-increase and multiplicative-decrease tracks a slowly varying neighbour and approaches the kink when $\tau_c \gg D$, but over-admits when $\tau_c \lesssim D$, since it keeps probing a signal that has become noise. A fixed budget is safe everywhere and forfeits $\Delta(D)$ when $\tau_c \gg D$. Which of the two is right depends on a quantity the controller can measure, because $r(D)$ is an autocorrelation of the probe stream the controller already collects.

Algorithm 1 estimates it online and switches. It holds one scalar of state beyond the AIMD window, costs constant time per epoch, and inherits Prop. 1, so it never exceeds the pod cap.

The tracking analysis of Algorithm 1 assumes (A1) stationary capacity $a_t \equiv a$ with $w_{\min} \leq a \leq K$, (A2) reliable detection of an overshoot within D epochs, and (A3) fixed $\alpha, \beta, w_{\min}, K$.

Proposition 3 (Sawtooth realization of the kink). *Under stationary capacity a and (A1) to (A3), the trackable branch of Algorithm 1 forms a periodic sawtooth about a with long-run over-admission $\rho = D\alpha / ((1-\beta)(a+D\alpha)) \approx D\alpha / ((1-\beta)a)$. Choosing $\alpha \leq \epsilon(1-\beta)a/D$ gives $\rho \leq \epsilon$, and $\rho \rightarrow 0$ as $D \rightarrow 0$.*

Proof sketch. After a backoff the window restarts at $\beta(a + D\alpha)$ and rises by α per clean epoch, overshooting a for D epochs before detection triggers the next backoff at peak $a + D\alpha$. The fraction of violating epochs per cycle gives ρ . This is the dynamic approach to the static-capacity kink of Thm. 1,

Algorithm 1 Adaptive admission for one contended unit

Require: cap K , floor $w_{\min}$, step α , backoff β , switch level θ , buffer length n , fixed budget $\bar{a}$

```

1:  $w \leftarrow w_{\min}$ 
2: for each control epoch  $t$  do
3:   append the newest probe reading to a buffer of the last
      $n$  readings
4:    $\hat{r} \leftarrow$  empirical correlation of readings  $D$  epochs apart
5:   if  $\hat{r} > \theta$  then ▷ neighbour is trackable
6:     if the reading from  $D$  epochs ago showed over-
       admission then
7:        $w \leftarrow \max(w_{\min}, \lceil \beta w \rceil)$ 
8:     else
9:        $w \leftarrow \min(w + \alpha, K)$ 
10:    end if
11:  else ▷ not trackable, so stop paying for the loop
12:     $w \leftarrow \min(\lfloor \bar{a} \rfloor, K)$ 
13:  end if
14:  admit  $w$  pods
15: end for
```

and D enters both here and in (3), which is the hinge between the systems measurement and the theory. $\square$

Theorem 2 (Near-optimality of Algorithm 1). *Let capacity be a symmetric two-state process with spectral gap $\gamma = 2 \min(p, 1-p)$. The empirical correlation is consistent, with $\Pr[|\hat{r}_n - r(D)| > t] \leq 2e^{-c_0 \gamma n t^2 / D^2}$, where the D^2 comes from the $2D$ -Lipschitz map $p \mapsto (1-2p)^D$ so the constant carries explicit D and γ dependence. Hence on any class with $\gamma \geq \gamma_0 > 0$, Algorithm 1 attains $\Delta(D)$ with regret*

Algorithm 2 Layered admission for several contended units

Require: cap K , per-unit switch levels θ_u , per-unit fixed decisions $\bar{m}_u$, buffer length n

```
1: for each control epoch  $t$  do
2:    $w \leftarrow 0$ 
3:   for  $u = 1$  to  $K$  do ▷ layers are nested, so
      $x_t^1 \geq x_t^2 \geq \dots$ 
4:      $\hat{r}_u \leftarrow$  empirical lag- $D$  correlation of layer  $u$ 
5:     if  $\hat{r}_u > \theta_u$  then
6:        $m_u \leftarrow 1$ [layer  $u$  was free  $D$  epochs ago]
7:     else
8:        $m_u \leftarrow \bar{m}_u$ 
9:     end if
10:     $w \leftarrow w + m_u$ 
11:  end for
12:  admit  $\min(w, K)$  pods
13: end for
```

$O(D/\sqrt{\gamma_0 N})$ over horizon N (App. B), and $O(1/N)$ for any fixed process. The restriction $\gamma \geq \gamma_0$ is needed, since as $\tau_c \rightarrow \infty$ we have $\gamma \rightarrow 0$ while sensing becomes trivial, so a γ -free uniform rate is neither attainable nor useful. The two separations are the endpoints of Cor. 1. Fixed AIMD's excess over-admission stays bounded below as $\tau_c/D \rightarrow 0$, and a fixed budget forfeits a constant of goodput as $\tau_c/D \rightarrow \infty$.

Proof sketch. The concentration bound gives a correct branch decision after $n = O(\log(1/\delta)/(r-\theta)^2)$ probes with probability $1-\delta$. Conditioned on the correct branch, Cor. 1 bounds the fixed-budget branch and Prop. 3 bounds the tracking branch, and averaging over the $O(1)$ misclassified prefix yields the regret. At small τ_c fixed AIMD sits near $\rho = \frac{1}{4}$ while the optimum tends to 0, and at large τ_c a fixed budget forfeits Δ tending to $\frac{1}{2}$. $\square$

Proposition 2 lifts admission past a single unit. Algorithm 2 runs one binary decision per contended layer and adds them, which is a valid scalar window because the layers are nested, and its advantage is bounded by the per-unit sum (4) and attains it under the monotonicity condition there. It reduces to Algorithm 1 when only one layer fluctuates. The condition is worth stating operationally. Equality needs $\pi_u(1 - r_u(D))$ non-increasing in u , so it holds when every layer shares one correlation, as it does for a birth-death capacity process with a height-independent flip rate, and it can fail when the higher layers are markedly less predictable than the lower ones. Algorithm 2 is correct and safe either way, since Prop. 1 applies per layer. What the condition governs is only whether the achieved advantage reaches the summed ceiling.

Scope. Three assumptions earn the closed form. First, per layer the unit is binary and symmetric. For any symmetric stationary x_t the value $\frac{1}{2}r(D)$ is achievable by the memoryless threshold $m_t = \hat{x}_t$, which uses only the freshest sample, and it is the exact advantage when x_t is Markov, since then $\hat{x}_t$ is a sufficient statistic. For a general non-Markov symmetric process a history-using controller can do strictly better, as

a random-phase square wave with $r(D)=0$ stays exactly predictable from its phase, so the equality is Markov-specific and the general converse is the mixing bound of App. A. Second, stationarity, since $r(D)$ is a lag- D autocorrelation of a stationary process. Under slow drift it is replaced by a local estimate, which is what Algorithm 1 already tracks. Third, fixed age D . If the freshest sample has a random age A the controller can read, the advantage is the age-averaged $\frac{1}{2} \mathbb{E}[r(A)]$, the same law under the age law, and the impossibility persists whenever $A \gg \tau_c$ almost surely.

VI. REALIZATION: SENSING OVER A FEDERATED FABRIC

The controller runs inside the cluster, behind the network boundary, and reaches the home node over a NATS [13] leaf link that needs one outbound port (Fig. 1). It probes cluster occupancy and actuates admission across that link. Submissions flow on `burst.submit` and results return on `results.>` by interest propagation, with no inbound cluster access. What matters for the controller is that this view is late. The leaf round trip and the CUDA start-up of a freshly admitted pod together set the sensing age of §IV, which we measure in §VIII. Tensors crossing the link are compressed by random-rotation quantization [14], [15], which helps bandwidth on the hop but does not reduce the sensing age, since that age is dominated by CUDA start-up rather than transport. The fabric's own design, its leaf federation, durability and single-port traversal, is characterized in the companion paper [2].

VII. IMPLEMENTATION

PoliteBurst is a Go control plane with a decider, prober, submitter, NATS bridge and back-off modules, plus a Python client and an IPython cell magic [16]. The controller runs inside the cluster under a minimal namespace-scoped ServiceAccount holding Job create, get and delete rights and pod read, connects to the in-cluster NATS leaf, and renders each `JobDescriptor` into a Kubernetes Job. Node targeting through a `nodeSelector` on the descriptor pins guest pods to non-reserved GPU types. We verified the full path live on the reference host [1]. A submit from the home node propagated through hub, leaf and controller, launched a GPU worker that ran real CUDA work, and published its result back through the leaf to the home driver. The system, the controller and worker images, and the transport and right-sizing packages [14], [17] are available as open-source artifacts, with repositories anonymized for review.

VIII. EVALUATION

We evaluate on a testbed whose sensing age is measured rather than assumed. A home node and a workstation-class node with GV100 GPUs run the controller, which probes real GPU occupancy and admits real CUDA workers while a competing tenant drives capacity. The fabric carrying this control plane is characterized in the companion paper [2]. Unless noted we report the mean and a t -based 95% confidence interval over independent repeats. Four questions organize the study. **(Q1)** What is the sensing age D that governs Prop. 3.

(Q2) Does Algorithm 1 stay polite while beating a fixed budget when it holds a self-imposed budget. (Q3) Does that hold under real time-varying contention, and where does it fail. (Q4) Does the value $\frac{1}{2}r(D)$ of Thm. 1 hold at the controller level, and does it also bind a learned controller.

A. Sensing age and the violation bound (Q1)

The sensing age has two measured parts. A read-only `nvidia-smi` probe costs 27 to 33 ms at the median. The end-to-end delay from launching a real GPU load until it is visible as utilization above the 40% floor is 2.46 s over 10 thermally guarded repeats on a GV100, with median 2.48 and range 2.29 to 2.61. CUDA start-up dominates that figure, not the poll. Which value applies depends on the question. Detecting whether a newly admitted pod is productive pays the full start-up cost, while probing pods that are already running costs only the poll. One epoch is a single completion-gated admission decision whose measured mean duration is about 1.1 s, so the 2.46 s start-up age is $D \approx 2$ to 3 epochs, which we sweep in Q4, and the 27 ms poll is below one epoch for running pods. Substituting into Prop. 3 makes D a first-order term in achievable politeness. A guest that admits aggressively with large α pays for it in over-admission scaled by this measured start-up delay.

B. Politeness under a self-imposed budget (Q2)

We burst B jobs under three policies against a self-imposed politeness budget C , measuring goodput in tasks per second and the over-budget fraction ρ , the share of epochs with concurrency above C . The policies are *naive*, which submits up to the hard cap, *static*, a fixed submit rate, and *aimd*, the trackable branch of Algorithm 1. To isolate the policy from the exogenous scheduling and image-pull variance of the shared cluster, which produced three pre-registered nulls in situ and is reported as such, we run on a controlled node with local-process workers and no scheduler or registry jitter. The design was pre-registered with randomized interleaved trial order and four repeats per cell, at two scales.

Table I reports it. Algorithm 1 holds the budget exactly at both scales, with $\rho = 0$, while delivering 1.72 times (GPU) and 3.18 times (CPU) the goodput of a fixed rate at the same politeness. At the GPU scale it also matches naive’s goodput without exceeding the budget. At the CPU scale, where naive over-admits flagrantly at $\rho = 0.95$, naive buys about twice the goodput and Algorithm 1 trades that headroom for politeness. Both pre-registered hypotheses hold with disjoint intervals at both scales. This isolates the admission rule on a controlled node, and the full pod-scheduling path is verified separately in §VII. Here the budget is self-imposed, so §VIII-C removes that assumption.

C. Politeness under real contention (Q3)

The harder question is politeness when capacity is genuinely scarce. We manufacture controlled contention on a two-GPU node with $C=2$. A competing tenant occupies $c(t) \in \{0, 1\}$ GPUs on a fixed profile, so the guest’s available capacity

TABLE I
GOODPUT AGAINST POLITENESS UNDER A SELF-IMPOSED BUDGET, 4 REPEATS, MEAN WITH 95% INTERVAL. ALGORITHM 1 HOLDS THE BUDGET AT A FIXED RATE’S POLITENESS WITH 1.7 TO 3.2 TIMES ITS GOODPUT.

scale	policy	goodput (tasks/s)	over-budget ρ
GPU, $C=2$	naive	0.0538 ± 0.0003	0.54 ± 0.01
	aimd	0.0539 ± 0.0001	0.00 ± 0.00
	static	0.0313 ± 0.0000	0.00 ± 0.00
CPU, $C=8$	naive	0.5779 ± 0.0062	0.95 ± 0.00
	aimd	0.2928 ± 0.0029	0.00 ± 0.00
	static	0.0920 ± 0.0002	0.00 ± 0.00

$a(t) = C - c(t)$ is exogenous, time-varying and measured by the prober through a per-GPU process census subject to the sensing age of Q1. Two profiles bracket the space. *Square* load is a period- 2τ wave modelling slowly varying or scheduled neighbours, and *poisson* load is memoryless on and off arrivals, the adversarial case. Alongside goodput and ρ we measure harm to the neighbour as the competitor’s completed work relative to its undisturbed baseline. The study was pre-registered with randomized interleaved order and four seeds per cell.

Table II reports a sharper picture than the self-imposed budget did. Under structured load Algorithm 1 is the polite and efficient knee. It reaches 1.41 times a fixed rate’s goodput, 0.091 against 0.064 tasks/s with disjoint intervals, at $\rho = 0.15$, which is a quarter of naive’s 0.64, and it spares the neighbour, retaining 76% of its throughput against naive’s 62%. The target ϵ and switch level θ were fixed before the runs. Host policy penalizes persistently under-utilized requests rather than momentary dips, so ϵ between 0.1 and 0.2 is the tolerable share of under-utilized epochs, and the measured 0.15 sits inside it. Naive’s higher goodput of 0.101 is bought entirely at the neighbour’s expense, since it over-admits two-thirds of the time and takes 38% of the competitor’s throughput.

Under memoryless load feedback breaks even and we report it. Goodput of 0.062 is statistically indistinguishable from the fixed rate’s 0.064 with overlapping intervals, and over-admission is worse at $\rho = 0.30$, because capacity now varies on the timescale of the sensing delay. By the time the prober sees a competitor arrive the observation is stale. This is a pre-registered null and it is what Thm. 1 predicts, since the two profiles are two coherence times on the same curve (Cor. 2). Slow square load has $\tau_c \gg D$ and gives feedback room to help, while memoryless load drives $r(D)$ toward zero, where no causal controller beats a fixed budget. The break-even is therefore a consequence of the theorem rather than a tuning failure, and it is the case Algorithm 1 is built to detect. The measured sensing path also places the Prop. 3 sawtooth ceiling above the observed ρ in both profiles, which is conservative for a discrete completion-gated controller, and its prediction that politeness degrades as contention outpaces D is borne out. Naive is the most impolite in every case, at ρ between 0.64 and 0.85.

TABLE II

POLITENESS UNDER REAL CONTENTION WITH $C=2, 4$ SEEDS, MEAN WITH 95% INTERVAL. “KEPT” IS COMPETITOR THROUGHPUT RETAINED, WHERE 1 IS UNDISTURBED. UNDER SQUARE LOAD ALGORITHM 1 IS THE POLITE AND EFFICIENT KNEE, AND UNDER POISSON LOAD FEEDBACK BREAKS EVEN, A PRE-REGISTERED NULL.

contention	policy	goodput (tasks/s)	ρ	kept
square	naive	0.101 ± 0.002	0.64 ± 0.01	0.62
	aimd	0.091 ± 0.000	0.15 ± 0.01	0.76
	static	0.064 ± 0.000	0.11 ± 0.01	0.89
poisson	naive	0.096 ± 0.002	0.85 ± 0.06	0.62
	aimd	0.062 ± 0.006	0.30 ± 0.06	0.55
	static	0.064 ± 0.000	0.22 ± 0.04	0.79

D. The value of feedback at the controller level (Q4)

The two contention profiles are two points, while Thm. 1 predicts the whole curve. To test it we simulate the window controllers on the two-state capacity process the theory analyzes, a Markov competitor with per-epoch flip probability p so the coherence time τ_c sweeps continuously, at the measured sensing age $D=3$ epochs with 1.2×10^5 epochs per point. Fig. 3 reports each controller’s goodput advantage over the best fixed budget at matched over-admission against the predicted $\frac{1}{2}r(D)$.

Four readings confirm the theory and motivate Algorithm 1. First, the threshold policy sits on the predicted line across three decades of τ_c , with maximum deviation 0.013, so the closed form is the exact frontier rather than an asymptote. Second, fixed AIMD captures roughly half the available advantage when $\tau_c \gg D$ and goes negative, meaning worse than a fixed budget, once $\tau_c \lesssim D$, because it keeps paying over-admission for a signal that has become noise. That is the same failure the live poisson case shows. Third, Algorithm 1, estimating $r(D)$ online, tracks the upper envelope. It never drops below a fixed budget and recovers most of AIMD’s advantage where feedback is worth having. Fourth, the value binds a learned controller. A policy-gradient agent in the DeepRM family [18], using a tabular softmax with a per-state baseline [19] and trained on the age- D observation history of four delayed readings, which is strictly more information than the memoryless threshold uses, converges to the threshold policy’s advantage within 0.003 at every point and exceeds $\frac{1}{2}r(D)$ by at most 0.013, the threshold policy’s own sampling deviation. Learning finds the frontier and does not pass it, which is the ceiling Thm. 1 sets for that family [18], [20]. Sweeping D collapses every curve onto $\frac{1}{2}e^{-D/\tau_c}$ against D/τ_c (Fig. 3, right), confirming the predicted scaling. The live square and poisson cases of Q3 are the two ends of that axis. **Live confirmation on GPU hardware.** We ran Algorithm 1 on the same two-GPU node across the Markov τ_c sweep, with p from 0.01 to 0.4, $D=3$ and two seeds, estimating $r(D)$ online. Two claims transfer from the model. The online estimate lands on the predicted curve, with $\hat{r} = (1 - 2\hat{p})^D$ tracking $(1 - 2p)^D$ to within 0.07 (Fig. 4, left). And the switch modulates as predicted, its closed-loop fraction falling




Fig. 3. Controller-level validation. *Left*. Goodput advantage over a fixed budget against capacity coherence time τ_c at $D=3$. The threshold policy lies on the predicted $\frac{1}{2}r(D)$, dashed, to within 0.013, and the DeepRM-style policy-gradient agent, diamonds, converges onto it without passing it. Fixed AIMD falls below a fixed budget once $\tau_c \lesssim D$, and Algorithm 1 tracks the envelope. *Right*. Sweeping D collapses the advantage onto $\frac{1}{2}e^{-D/\tau_c}$ against D/τ_c .




Fig. 4. Algorithm 1 live on two GV100 GPUs, Markov τ_c sweep, $D=3$, two seeds. *Left*. The online estimate $\hat{r}$ and the realized availability autocorrelation land on $(1 - 2p)^D$, and the switch fires as $r(D)$ crosses θ . *Right*. The goodput and over-admission frontier, where the adaptive controller tracks AIMD when feedback is useful and moves toward a fixed budget’s politeness as p grows.

from 0.98 to 0.48 as $r(D)$ drops past θ . Hardware differs from simulation in one instructive way. Raw goodput rewards over-admission, so live AIMD out-performs a fixed budget everywhere, and at high flip rate it buys that goodput with impoliteness at $\rho = 0.44$ against 0.04 rather than by tracking. The predicted crossover therefore does not appear in raw goodput. It appears on the goodput and politeness frontier (Fig. 4, right) as the over-admission the switch sheds against always running AIMD, which grows from +0.01 to +0.22 as p rises from 0.08 to 0.4. Algorithm 1 matches AIMD where the neighbour is trackable and recovers a fixed budget’s politeness where it is not.

E. Summary and threats to validity

The evaluation establishes four things. The sensing age is a measured first-order term at about 2.5s, dominated by CUDA start-up. Under a self-imposed budget Algorithm 1 delivers 1.7 to 3.2 times a fixed rate’s goodput at $\rho = 0$. Under structured contention it is the polite and efficient knee and spares the neighbour, while under memoryless contention feedback breaks even, a pre-registered null we report rather than suppress. And the value $\frac{1}{2}r(D)$ holds at the controller level, with the threshold policy on it to within 0.013, a learned policy-gradient agent capped by it, and Algorithm 1 never below a fixed budget, which unifies the two live contention cases as endpoints of one curve.

Three threats remain. Contention is manufactured on a two-GPU node at $C=2$ with synthetic profiles rather than drawn from a production multi-tenant trace. The clean tracing of the advantage in Q4 is controller-level, because live static is a fixed trickle with no tunable frontier to measure an advantage against. And the live adaptive sweep uses two seeds per point, so it is a confirmatory anchor rather than primary evidence. The full system path is nonetheless verified end-to-end in §VII, and naive’s impoliteness, at $\rho \geq 0.64$ with at least 38% neighbour harm, is robust across every case.

IX. RELATED WORK

Bursting and function execution. HTCondor [3], funcX and Globus Compute [4], Parsl [5], Ray [6] and Dask [7] provide elastic task execution and provider-driven scaling, and treat the host as owned capacity to fill. None of them models a fair-use politeness budget or admission under a delayed view of capacity. They are complementary substrates, and a head-to-head comparison against them appears in the companion systems paper [2].

Congestion and scavenger control. Our loop is AIMD in the lineage of additive-increase and multiplicative-decrease congestion avoidance [21], and its objective of yielding to incumbents and using only slack aligns with low-priority transports such as LEDBAT [22] and TCP Nice [23]. Those protocols back off on one-sided loss or delay signals. Two things differ here. The objective is two-sided, since the guest must stay under a pod cap and above a utilization floor, which admits an over-admission measure rather than pure loss. And we bound what any such rule can achieve rather than proposing one more back-off schedule.

Learned resource control. Deep reinforcement learning schedulers including DeepRM [18] and Decima [20], with variance reduction for exogenous input processes [19], learn scheduling policies from cluster feedback. Our ceiling applies to them unchanged, since it constrains any causal policy rather than any particular family, and §VIII-D tests a policy-gradient agent of that family against it.

Age of information. The ceiling is a statement about freshness, since the value of feedback is set by the age of the capacity estimate relative to how fast the neighbour changes. That places polite bursting in the age-of-information line of work [24], on the control side rather than the sampling side. We bound the achievable goodput and politeness frontier as a function of sensing age and give algorithms that adapt to the measured value.

X. CONCLUSION

Restraint on a shared cluster is a control problem whose difficulty is fixed by how stale the guest’s view of the cluster is. We gave that dependence an exact form. The extra goodput a feedback controller can win over the best fixed budget is half the correlation between capacity now and capacity as last seen, which is attained by admitting work whenever the last reading showed capacity free, and which vanishes once the sensing delay outgrows the neighbour’s correlation time.

Two algorithms follow from that form. One measures the correlation online and switches between feedback and a fixed budget, staying at or above the better of the two. The other runs a binary decision per contended unit and so covers capacities beyond a single unit. Both carry the safety property that the pod cap is never exceeded, which holds deterministically rather than in expectation.

Short controlled experiments support the result on both sides of the crossover. Under structured neighbour load, feedback admission reaches 1.41 times a fixed budget’s goodput at a quarter of a greedy policy’s over-admission while leaving the competing tenant more throughput. Under memoryless load it breaks even, and the switching algorithm detects the case and stops paying for a loop that cannot help. A learned policy-gradient controller given a longer observation history converges to the ceiling and does not pass it.

Three things remain open. Our contention is manufactured on a two-GPU node rather than drawn from a production multi-tenant trace. The exact equality holds for a binary Markov unit, while the general statement is the weaker mixing bound of App. A. And a controller that estimates the correlation itself, rather than only which side of the switch it is on, should recover part of the remaining distance to the ceiling.

APPENDIX A

GENERAL-CAPACITY IMPOSSIBILITY VIA β -MIXING

Corollary 1 showed that feedback’s value vanishes with the sensing age D for the two-state model. We now show the same for any stationary capacity process, with the decay rate governed by its β -mixing coefficient, also called absolute regularity. The exact binary law is one tight instance of that general statement.

Setup. Let $\{a_t\}$ be stationary on $\{0, \dots, C\}$ with $\mathcal{F}_{\leq s} = \sigma(a_u : u \leq s)$ and $\mathcal{F}_{\geq s} = \sigma(a_u : u \geq s)$. Its β -mixing coefficient at lag D is

$$\beta(D) = \mathbb{E} \left[\sup_{B \in \mathcal{F}_{\geq t}} |\mathbb{P}(B \mid \mathcal{F}_{\leq t-D}) - \mathbb{P}(B)| \right], \quad (5)$$

and $\beta(D) \rightarrow 0$ holds if and only if the process is absolutely regular. Geometrically mixing processes satisfy $\beta(D) \leq c_0 \rho^D$ for some $\rho \in (0, 1)$. Note that $\beta(D)$ always carries its lag argument and is unrelated to the AIMD back-off factor β . Fix any bounded per-epoch objective $R(w, a)$, for instance the politeness Lagrangian $R_\lambda(w, a) = \min(w, a) - \lambda(w - a)^+$, with oscillation $\Omega \triangleq \max_{w,a} R - \min_{w,a} R$. A causal age- D controller sets w_t measurable with respect to $\mathcal{G} \triangleq \mathcal{F}_{\leq t-D}$, while a fixed controller uses a $\mathcal{G}$ -independent w . Write $V^*(D)$ and V_{stat} for their optimal per-epoch values and $\Delta(D) = V^*(D) - V_{\text{stat}}$.

Lemma 2 (Lipschitz value function). $\psi(\mu) \triangleq \max_w \sum_a R(w, a) \mu(a)$ is convex on the simplex and Ω -Lipschitz in total variation, so $|\psi(\mu) - \psi(\mu')| \leq \Omega \|\mu - \mu'\|_{\text{TV}}$. Theorem 3 below is then the standard value-of-information bound [10], [11] specialized to the mixing coefficient.

Proof. ψ is a pointwise maximum of linear functionals, hence convex. Let w_μ attain $\psi(\mu)$. Then $\psi(\mu) - \psi(\mu') \leq \sum_a R(w_\mu, a)(\mu - \mu')(a) \leq \text{osc}_a R(w_\mu, \cdot) \|\mu - \mu'\|_{\text{TV}} \leq \Omega \|\mu - \mu'\|_{\text{TV}}$, using $\sum_a f(a)\nu(a) \leq \text{osc}(f) \|\nu\|_{\text{TV}}$ for zero-sum ν . The reverse bound is symmetric. $\square$

Theorem 3 (β -mixing impossibility). *For any stationary capacity process and any bounded objective,*

$$0 \leq \Delta(D) \leq \Omega \beta(D). \quad (6)$$

Hence no causal age- D controller outperforms a fixed budget as $D \rightarrow \infty$ whenever the capacity is absolutely regular. If $\beta(D) \leq c_0 \varrho^D$ the advantage decays geometrically, with $\Delta(D) \leq \Omega c_0 \varrho^D$.

Proof. Let $\mu_{\mathcal{G}} = \mathbb{P}(a_t \in \cdot \mid \mathcal{G})$, so the tower property gives $\mathbb{E}[\mu_{\mathcal{G}}] = \bar{\mu} \triangleq \mathbb{P}(a_t \in \cdot)$. Optimizing the $\mathcal{G}$ -measurable window pointwise, $V^*(D) = \mathbb{E}[\psi(\mu_{\mathcal{G}})]$ and $V_{\text{stat}} = \psi(\bar{\mu})$, so convexity of ψ and Jensen give $\Delta(D) = \mathbb{E}[\psi(\mu_{\mathcal{G}})] - \psi(\bar{\mu}) \geq 0$, meaning fresher information never hurts. For the upper bound, Lemma 2 gives $\Delta(D) \leq \Omega \mathbb{E} \|\mu_{\mathcal{G}} - \bar{\mu}\|_{\text{TV}} = \Omega \mathbb{E}[\sup_B |\mathbb{P}(a_t \in B \mid \mathcal{G}) - \mathbb{P}(a_t \in B)|] \leq \Omega \beta(D)$. The last step holds because $\sigma(a_t) \subseteq \mathcal{F}_{\geq t}$, so (5) takes the supremum over a larger field. $\square$

Corollary 3 (Binary tightness). *For the symmetric two-state model, $\beta(D) = \frac{1}{2}|1 - 2q(D)| = \frac{1}{2}r(D)$ by a standard mixing identity [12], [8], so Thm. 3 reproduces the exact advantage of Thm. 1 up to the objective scale Ω .*

Proof. The Markov property reduces $\mathcal{F}_{\leq t-D}$ to a_{t-D} . The conditional-to-marginal total-variation distance is $|\frac{1}{2} - q(D)|$ from either state, so $\beta(D) = \frac{1}{2} - q(D) = \frac{1}{2}r(D)$. $\square$

APPENDIX B

REGRET OF THE ADAPTIVE CONTROLLER

We sketch the near-optimality asserted in Thm. 2. Algorithm 1 estimates $\hat{p}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{1}[\hat{a}_i \neq \hat{a}_{i-1}]$, forms $\hat{r}_n = (1 - 2\hat{p}_n)^D$, and runs the closed-loop branch when $\hat{r}_n > \theta$. The flip indicators are a bounded functional of the reversible two-state chain, whose absolute spectral gap is $\gamma = 1 - |1 - 2p| = 2 \min(p, 1 - p)$, so the Hoeffding inequality for reversible chains [25], [26] gives $\Pr(|\hat{p}_n - p| \geq t) \leq 2e^{-c_0 \gamma n t^2}$ for a universal $c_0 > 0$. Since $p \mapsto (1 - 2p)^D$ is $2D$ -Lipschitz on $[0, \frac{1}{2}]$, we have $|\hat{r}_n - r(D)| \leq 2D|\hat{p}_n - p|$. With a threshold margin $|r(D) - \theta| \geq \Delta\theta$ the branch is therefore wrong with probability at most $2 \exp(-c_0 \gamma n (\Delta\theta)^2 / 4D^2)$.

Theorem 4 (Regret). *Set θ at the indifference level, where the two branches' values are L -Lipschitz in $r(D)$ and coincide. On any process class with spectral gap $\gamma \geq \gamma_0 > 0$, the regret of Algorithm 1 over horizon N is $O(D/\sqrt{\gamma_0 N})$, and $O(1/N)$ for any fixed process with $r(D) \neq \theta$.*

Proof sketch. A wrong branch costs at most $L\Delta\theta$. Averaging the wrong-branch bound gives $\text{Regret}_N \leq L \min(\Delta\theta, 8D^2/(c_0 \gamma N \Delta\theta))$, optimized at $\Delta\theta^* = \sqrt{8D^2/(c_0 \gamma N)}$ to $O(D/\sqrt{\gamma N})$, and $O(1/N)$ for fixed $\Delta\theta$. $\square$

REFERENCES

- [1] D. Weitzel, A. Graves, S. Albin, H. Zhu, F. Würthwein, M. Tatineni, D. Mishin, J. Graham, E. E. Khoda, M. F. Sada, L. Smarr, and T. DeFanti, "The national research platform: Stretched, multi-tenant, scientific Kubernetes cluster," *arXiv preprint arXiv:2505.22864*, 2025.
- [2] Anonymous Author(s), "A policy-aware, container-native bursting fabric with federated messaging, transport compression, and thermal-aware right-sizing," Companion systems paper; anonymized for double-blind review, available to reviewers on request.
- [3] D. Thain, T. Tannenbaum, and M. Livny, "Distributed computing in practice: the Condor experience," *Concurrency and Computation: Practice and Experience*, vol. 17, no. 2-4, pp. 323-356, 2005.
- [4] R. Chard, Y. Babuji, Z. Li, T. Skluzacek, A. Woodard, B. Blaiszik, I. Foster, and K. Chard, "funcX: A federated function serving fabric for science," in *Proc. 29th Int. Symp. on High-Performance Parallel and Distributed Computing (HPDC)*, 2020.
- [5] Y. Babuji, A. Woodard, Z. Li *et al.*, "Parsl: Pervasive parallel programming in Python," in *Proc. 28th Int. Symp. on High-Performance Parallel and Distributed Computing (HPDC)*, 2019.
- [6] P. Moritz, R. Nishihara, S. Wang *et al.*, "Ray: A distributed framework for emerging AI applications," in *Proc. 13th USENIX Symp. on Operating Systems Design and Implementation (OSDI)*, 2018.
- [7] M. Rocklin, "Dask: Parallel computation with blocked algorithms and task scheduling," in *Proc. 14th Python in Science Conference (SciPy)*, 2015.
- [8] D. A. Levin, Y. Peres, and E. L. Wilmer, *Markov Chains and Mixing Times*, 2nd ed. American Mathematical Society, 2017.
- [9] S. H. A. Ahmad, M. Liu, T. Javidi, Q. Zhao, and B. Krishnamachari, "Optimality of myopic sensing in multichannel opportunistic access," *IEEE Transactions on Information Theory*, vol. 55, no. 9, pp. 4040-4050, 2009.
- [10] R. A. Howard, "Information value theory," *IEEE Transactions on Systems Science and Cybernetics*, vol. 2, no. 1, pp. 22-26, 1966.
- [11] M. De Lara and O. Gossner, "Payoffs-beliefs duality and the value of information," *SIAM Journal on Optimization*, vol. 30, no. 1, pp. 464-489, 2020.
- [12] R. C. Bradley, "Basic properties of strong mixing conditions: A survey and some open questions," *Probability Surveys*, vol. 2, pp. 107-144, 2005.
- [13] "NATS: Connective technology for adaptive edge and distributed systems," <https://nats.io>, accessed 2026.
- [14] Anonymous Author(s), "A random-rotation quantization library for embedding and KV-cache compression," Software artifact; repository anonymized for double-blind review, available to reviewers on request.
- [15] A. Zandieh, M. Daliri, M. Hadian, and V. Mirrokni, "TurboQuant: Online vector quantization with near-optimal distortion rate," *arXiv preprint arXiv:2504.19874*, 2025.
- [16] Anonymous Author(s), "The bursting fabric described in this paper," Software artifact; repository anonymized for double-blind review, available to reviewers on request.
- [17] —, "An OOM-safe GPU batch-sizing and thermal-aware control library," Software artifact; repository anonymized for double-blind review, available to reviewers on request.
- [18] H. Mao, M. Alizadeh, I. Menache, and S. Kandula, "Resource management with deep reinforcement learning," in *Proc. 15th ACM Workshop on Hot Topics in Networks (HotNets)*, 2016.
- [19] H. Mao, S. B. Venkatakrishnan, M. Schwarzkopf, and M. Alizadeh, "Variance reduction for reinforcement learning in input-driven environments," in *Proc. 7th Int. Conf. on Learning Representations (ICLR)*, 2019.
- [20] H. Mao, M. Schwarzkopf, S. B. Venkatakrishnan, Z. Meng, and M. Alizadeh, "Learning scheduling algorithms for data processing clusters," in *Proc. ACM SIGCOMM*, 2019.
- [21] D.-M. Chiu and R. Jain, "Analysis of the increase and decrease algorithms for congestion avoidance in computer networks," *Computer Networks and ISDN Systems*, vol. 17, no. 1, pp. 1-14, 1989.
- [22] S. Shalunov, G. Hazel, J. Iyengar, and M. Kuehlewind, "Low extra delay background transport (LEDBAT)," RFC 6817, IETF, 2012.
- [23] A. Venkataramani, R. Kokku, and M. Dahlin, "TCP Nice: A mechanism for background transfers," in *Proc. 5th USENIX Symp. on Operating Systems Design and Implementation (OSDI)*, 2002.

- [24] R. D. Yates, Y. Sun, D. R. Brown, S. K. Kaul, E. Modiano, and S. Ulukus, "Age of information: An introduction and survey," *IEEE Journal on Selected Areas in Communications*, vol. 39, no. 5, pp. 1183–1210, 2021.
- [25] D. Paulin, "Concentration inequalities for Markov chains by Marton couplings and spectral methods," *Electronic Journal of Probability*, vol. 20, no. 79, pp. 1–32, 2015.
- [26] P. Lezaud, "Chernoff-type bound for finite Markov chains," *The Annals of Applied Probability*, vol. 8, no. 3, pp. 849–867, 1998.